

```

import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np
import matplotlib.pyplot as plt

# ===== Załaduj dane MNIST =====
(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.mnist.load_data()
x_train = x_train.astype("float32") / 255.
x_test = x_test.astype("float32") / 255.
x_train = x_train.reshape(-1, 28 * 28)
x_test = x_test.reshape(-1, 28 * 28)

# ===== Filtruj tylko klasy 0-4 do treningu =====
train_mask = y_train <= 4
x_train_filtered = x_train[train_mask]

# ===== Budowa autoenkodera =====
input_dim = 784
encoding_dim = 64

input_img = tf.keras.Input(shape=(input_dim,))
encoded = layers.Dense(128, activation='relu')(input_img)
encoded = layers.Dense(encoding_dim, activation='relu')(encoded)

decoded = layers.Dense(128, activation='relu')(encoded)
decoded = layers.Dense(input_dim, activation='sigmoid')(decoded)

autoencoder = models.Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='mse')
autoencoder.summary()

# ===== Trening autoenkodera =====
autoencoder.fit(x_train_filtered, x_train_filtered,
                epochs=20,
                batch_size=256,
                shuffle=True,
                validation_split=0.2)

# ===== Ewaluacja na pełnym zbiorze testowym (klasy 0-9) =====
reconstructions = autoencoder.predict(x_test)
reconstruction_errors = np.mean(np.square(x_test - reconstructions),
axis=1)

# ===== Ustal próg detekcji na podstawie klas 0-4 =====
test_mask_0_4 = y_test <= 4
threshold = np.mean(reconstruction_errors[test_mask_0_4]) + 2 *
np.std(reconstruction_errors[test_mask_0_4])

# ===== Wykrywanie anomalii =====

```

```

y_pred_anomaly = reconstruction_errors > threshold
y_true_anomaly = y_test > 4 # klasy 5-9 to anomalie

# ===== Ocena skuteczności =====
from sklearn.metrics import classification_report, confusion_matrix

print("Raport klasyfikacji (anomalie = 1):")
print(classification_report(y_true_anomaly, y_pred_anomaly))
print("Macierz pomyłek:")
print(confusion_matrix(y_true_anomaly, y_pred_anomaly))

# ===== Przykładowe błędy rekonstrukcji =====
plt.hist(reconstruction_errors[test_mask_0_4], bins=50, alpha=0.6,
label='Klasy 0-4')
plt.hist(reconstruction_errors[~test_mask_0_4], bins=50, alpha=0.6,
label='Klasy 5-9 (anomalie)')
plt.axvline(threshold, color='red', linestyle='--', label='Próg
detekcji')
plt.xlabel("Błąd rekonstrukcji")
plt.ylabel("Liczba próbek")
plt.legend()
plt.title("Histogram błędu rekonstrukcji")
plt.show()

```

Model: "functional"

Layer (type) Param #	Output Shape	
input_layer (InputLayer)	(None, 784)	
dense (Dense)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dense_2 (Dense)	(None, 128)	
dense_3 (Dense)	(None, 784)	


101,136 |

Total params: 218,192 (852.31 KB)

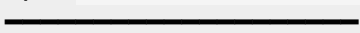
Trainable params: 218,192 (852.31 KB)

Non-trainable params: 0 (0.00 B)

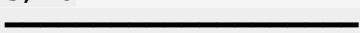
Epoch 1/20

96/96  3s 9ms/step - loss: 0.1210 - val\_loss: 0.0446

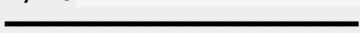
Epoch 2/20

96/96  1s 7ms/step - loss: 0.0407 - val\_loss: 0.0307

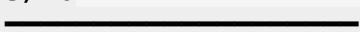
Epoch 3/20

96/96  1s 7ms/step - loss: 0.0288 - val\_loss: 0.0237

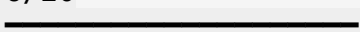
Epoch 4/20

96/96  1s 7ms/step - loss: 0.0226 - val\_loss: 0.0190

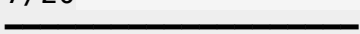
Epoch 5/20

96/96  1s 7ms/step - loss: 0.0182 - val\_loss: 0.0160

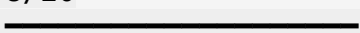
Epoch 6/20

96/96  1s 7ms/step - loss: 0.0154 - val\_loss: 0.0145

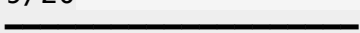
Epoch 7/20

96/96  1s 7ms/step - loss: 0.0139 - val\_loss: 0.0131

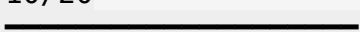
Epoch 8/20

96/96  1s 7ms/step - loss: 0.0129 - val\_loss: 0.0123

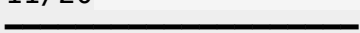
Epoch 9/20

96/96  1s 8ms/step - loss: 0.0118 - val\_loss: 0.0115

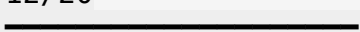
Epoch 10/20

96/96  1s 7ms/step - loss: 0.0112 - val\_loss: 0.0109


Epoch 11/20

96/96  1s 7ms/step - loss: 0.0106 - val\_loss: 0.0103

Epoch 12/20

96/96  1s 7ms/step - loss: 0.0100 - val\_loss: 0.0099

Epoch 13/20

96/96  1s 6ms/step - loss: 0.0096 - val\_loss: 0.0095

Epoch 14/20

96/96 \_\_\_\_\_ 1s 8ms/step - loss: 0.0092 - val\_loss: 0.0092

Epoch 15/20

96/96 \_\_\_\_\_ 1s 7ms/step - loss: 0.0089 - val\_loss: 0.0089

Epoch 16/20

96/96 \_\_\_\_\_ 1s 7ms/step - loss: 0.0086 - val\_loss: 0.0085

Epoch 17/20

96/96 \_\_\_\_\_ 1s 7ms/step - loss: 0.0083 - val\_loss: 0.0086

Epoch 18/20

96/96 \_\_\_\_\_ 1s 7ms/step - loss: 0.0081 - val\_loss: 0.0083

Epoch 19/20

96/96 \_\_\_\_\_ 1s 8ms/step - loss: 0.0079 - val\_loss: 0.0079

Epoch 20/20

96/96 \_\_\_\_\_ 1s 7ms/step - loss: 0.0077 - val\_loss: 0.0078

313/313 \_\_\_\_\_ 1s 2ms/step

Raport klasyfikacji (anomalia = 1):

	precision	recall	f1-score	support
False	0.58	0.96	0.72	5139
True	0.86	0.28	0.42	4861
accuracy			0.63	10000
macro avg	0.72	0.62	0.57	10000
weighted avg	0.72	0.63	0.58	10000

Macierz pomyłek:

[[4913 226]

[3503 1358]]

